

Agent Architecture

QuickBite Agentic Analytics

Natural-language business analytics answered by four cooperating agents that plan the analysis, write and execute SQL, verify the numbers arithmetically, and explain the result.

Repository: <https://github.com/anilgehltn/quickbite-agentic-analytics>

Source: docs/ARCHITECTURE.md — regenerate with scripts/build_architecture_pdf.py

1. Overview

QuickBite Agentic Analytics answers natural-language questions about a quick-service restaurant chain's sales data. A user types "which stores are declining and why"; the system plans an analysis, writes and executes SQL, checks the resulting numbers arithmetically, and explains what they mean — and shows every one of those steps, including the exact SQL, in a trace beside the answer.

The problem it solves is narrower than "natural language to SQL". Producing a query that runs is easy. Producing an answer that is *correct*, and being able to demonstrate that it is correct, is the hard part. A query can execute perfectly and return the wrong rows; a summary can quote accurate figures and draw a conclusion they do not support. Both failures look exactly like success.

One principle governs every design decision in this system:

The model decides what to compute. Deterministic code decides how it executes, and checks the result.


Concretely, agents emit structured artefacts — a plan conforming to a Pydantic schema, a SQL string, a JSON insight object. They never perform arithmetic, never format a number, never decide whether a value is plausible, and never determine which rows satisfy a condition. Everything a language model is unreliable at is done in Python or SQL; everything it is good at — interpreting intent, choosing an approach, writing prose — is left to it.

This principle is not stated once and forgotten. Sections 6 and 10 describe specific defects that occurred when it was violated, and the structural changes made in response.

Stack. Python 3.12, FastAPI, Pydantic v2, SQLite, sqlglot for SQL parsing. Next.js 14 with TypeScript and Tailwind for the frontend. Four LLM providers behind one client. No agent framework — see section 10 for why.

2. Agent architecture

Diagram



Text fallback

```

User question
|
v
Entity resolver (deterministic; resolves "ST7" -> ST007, "Bangalore" -> Bengaluru)
|
v
Planner agent <..... Semantic layer (schema, 10 metrics, 10 business rules)
|
| AnalysisPlan
v
[ intent? ] --- unsupported --> Scope reply + adjacent answerable questions --.
| --- ambiguous --> Clarifying question + interpretations -----|
| answerable |
v |
SQL analyst agent <.... Semantic layer |
| |
| SQL |
v |
SQL guard (single statement, query-type node, table allowlist, LIMIT) |
| |
v |
SQLite, read-only connection ---- rows ----> Verifier agent |
| |
13 deterministic checks |
| |
failed? --yes, once--> back to analyst |
| |
no |
v |
Insight agent |
| |
v v
Response: answer, chart, verification, trace

```

The semantic layer and the entity resolver are inputs to the agents, not stages in the pipeline. The SQL guard sits between the analyst and the database and cannot be bypassed: the analyst has no other route to SQLite.

The agents

Agent	Role	Input	Output	On failure
Planner	Interpret the question; decide what to measure, over what window, broken into which sub-queries	Question text, resolved entities, semantic layer	AnalysisPlan (intent, time window, metrics, dimensions, sub-queries, reasoning, confidence)	No deterministic substitute exists. The run stops and returns the unavailable message with the cached questions offered. One retry first, with the schema validation error appended.
SQL analyst	Write and execute one SQL query per sub-query	SubQuery + AnalysisPlan + semantic layer	QueryResult (SQL, columns, rows, row count, timing, attempts)	Two attempts, the second carrying the guard's or the database's own error message. If the <i>provider</i> is unreachable, SQL is assembled deterministically from the plan and the result is flagged degraded . If the provider is reachable but its SQL is rejected, the result carries the error — see section 10.
Verifier	Check the numbers against each other and against the plan	AnalysisPlan + all QueryResults	VerificationReport (status, checks, summary)	Every check is deterministic Python; a provider outage costs nothing. The optional model escalation reports that it could not run rather than being hidden.
Insight	Explain what the result means, in business language	AnalysisPlan + verified results	Insight (headline, narrative, key findings, caveats, actions, confidence) + ChartSpec	Falls back to deterministic prose generated from the rows, at reduced confidence, and the trace says the explanation layer was unavailable.

Every agent inherits from a common base that owns timing, structured logging, error capture, and production of the `AgentStep` that appears in the trace. A subclass implements only `execute` and `summarize`. That split is deliberate: if producing a trace step were each agent's own responsibility, the one agent that forgot would be invisible in the UI, and a failing agent would be invisible exactly when the trace is most useful.

3. Orchestration

The orchestrator sequences the agents and owns every degradation decision. The agents themselves are deliberately narrow — the planner does not know the analyst can fail, the analyst does not know what the verifier will make of its rows — so the policy lives in one readable file rather than scattered across four.

Sequence.

- 1 **Plan.** One planner call, under a 45-second timeout.
- 2 **Branch on intent.** `unsupported` returns a scope reply; `ambiguous` returns a clarifying question. Neither runs any SQL.
- 3 **Execute.** All sub-queries run **concurrently** via `asyncio.gather`, under a 90-second timeout for the batch. They are independent by construction, so on a diagnostic question with five sub-queries this is the difference between five sequential model round trips and roughly one.
- 4 **Verify**, under a 45-second timeout.
- 5 **Self-heal, at most once.** If verification fails on an error-severity check, only the sub-queries the report blames are re-run, with the verifier's own message appended to the analyst's prompt, and the merged results are

verified again. Exactly once: a model that cannot fix a stated arithmetic inconsistency on the second attempt will not fix it on the fifth, and a user is waiting. A repaired result never overwrites one that at least produced rows.

6 **Explain**, under a 90-second timeout.

Retries, at three levels, for different failures.

Level	Count	Trigger
HTTP, within a provider	2 retries, exponential backoff from 0.5 s	429, 5xx, timeout, connection reset
Provider failover	Every configured provider in order	Any provider-level failure
Planner schema	1 retry	Plan fails Pydantic validation
SQL	2 attempts total	Guard rejection or execution error
Verification self-heal	1	Error-severity check failed

Degradation policy. Each stage has a defined answer to "what if this fails", and the tests in `tests/test_degraded.py` simulate a dead provider at each stage independently.

- Planner fails → unavailable message; the eight cached questions are offered.
- Analyst fails on *some* sub-queries → answer proceeds on the rest, and the gap is visible in the trace and the results list.
- Analyst fails on *all* sub-queries → honest failure naming each reason.
- Verifier fails → the answer is still returned, marked unverified.
- Insight fails → deterministic prose from the rows, reduced confidence.

The pipeline never raises. Every path out of `Orchestrator.run` returns a valid `AnalysisResponse`. `run` wraps `_run` in a catch-all that converts an unexpected exception into a well-formed unanswered response with whatever trace was accumulated. A caller renders one shape whether the run succeeded, partly succeeded, timed out, or was refused; an exception can never reach a user as a 500.

The trace is complete and honest. Every stage contributes an `AgentStep`, including skipped and failed ones. A skipped verifier appears as `SKIPPED` with its reason, not as an absence. During a live run the frontend shows the pipeline advancing, but those in-flight steps deliberately carry no duration and no token count — the ordering is a schedule, not a measurement, and inventing numbers in the one component whose job is to be trustworthy would be self-defeating.

4. Data architecture

A star schema in SQLite, built from an Excel workbook by a deterministic ETL.

Table	Rows	Purpose
<code>fact_orders</code>	20,000	Order grain. Revenue, order counts, AOV, channel, hour, day type, festive flags
<code>fact_order_lines</code>	49,834	Line grain. SKU, quantity, line value, estimated COGS, line margin

Table	Rows	Purpose
dim_store	50	Store, city, region, format, price index
dim_product	30	SKU, category, veg/non-veg, base price, COGS percentage
dim_customer	5,000	Segment, home city, join date
dim_promotion	6	Promotion type, discount, applicability
dim_calendar	365	Day type, festive period, month key
mart_store_month	600	Pre-aggregated monthly store metrics
mart_city_month	96	Pre-aggregated monthly city metrics
mart_channel_month	48	Pre-aggregated monthly channel metrics

The three marts exist because monthly trend, ranking and decline questions are the most common shapes asked and the most expensive to compute repeatedly from 20,000 orders. They reconcile exactly with `fact_orders`, which the quality gate asserts.

Indexes. Eleven explicit indexes: `fact_orders` on order date, store, channel, month key, weekend flag and festive period; `fact_order_lines` on order id and SKU; each mart on month key.

Quality gate. `python -m app.etl.quality_checks` runs 36 checks over the built database — row counts, referential integrity, arithmetic reconciliation between grains, null profiles, date ranges, and the documented distributional properties. It currently reports **0 errors, 1 warning, 5 observations**. It runs in CI on a database rebuilt from the workbook on every push, so an ETL that stops reproducing the shipped data fails there rather than in a demo.

Three data decisions that shape everything

1. The data anchor is fixed, not read from the clock. The dataset covers 2025-08-01 to 2026-07-31. "Today" is the constant `DATA_ASOF_DATE = 2026-07-31` everywhere — in the planner's prompt, in the SQL rules, in the verifier's range check, and in the frontend footer. Any query using `DATE( 'now' )` returns zero rows against this data, so the SQL rules forbid clock-reading functions outright and the verifier independently asserts that every date in every result falls inside the data range. Without this, "last 3 months" silently means something different on every day the system is run, and a reviewer opening the deployment months later would get empty answers with no visible cause.

2. Revenue is tax-exclusive, and that is stated once. `net_before_tax` is canonical; `net_revenue` includes the 5% tax and overstates revenue by exactly that. The semantic layer defines `revenue` as `SUM(net_before_tax)` and a business rule states the distinction in the words an agent will encounter. Every answer's caveats repeat it. The alternative — letting each query decide — produces a set of answers that are each defensible and mutually inconsistent.

3. Order grain is canonical for revenue. The two grains do not reconcile exactly: line-level revenue exceeds order-level by **0.110% across the full year, but 0.445% within the May–July evaluation window** — the variance is concentrated in exactly the period the evaluation questions ask about. A revenue figure taken from `fact_order_lines` would therefore disagree with the canonical answer by an amount large enough to notice and small enough to look like a rounding artefact. So `fact_orders` is canonical for revenue, orders and AOV; `fact_order_lines` is used only for SKU, category and margin questions, where it is the only grain that has the data. This is a business rule in the semantic layer, not a convention.

5. The semantic layer

`app/semantic/schema.py` is a single module rendering roughly 17,600 characters of context that both the planner and the analyst receive. It contains:

- **The time anchor**, with "last 3 months" and similar phrases pre-resolved to absolute dates.
- **Every table and column**, with types and a one-line description of what the column actually means.
- **Ten canonical metric definitions**, each with its exact SQL expression, its source table and its unit — `revenue`, `revenue_with_tax`, `gross_revenue`, `discount`, `orders`, `units`, `aov`, `units_per_order`, `gross_margin`, `margin_pct`. A plan naming any other metric is rejected by a Pydantic validator before SQL is written.
- **Ten business rules** stating the traps in the words an agent will meet them: `LEFT JOIN dim_customer` because 28.32% of orders are anonymous walk-ins; `LEFT JOIN dim_promotion` because only 4.20% carry a promotion; `revenue` means `net_before_tax`; averages are ratios of sums, never `AVG()` over pre-aggregated rows.
- **Four few-shot query examples**, each executed against the real database by `tests/test_semantic.py`. Few-shot SQL referencing a column that does not exist teaches an agent to hallucinate the same column, so the examples are verified rather than trusted.

Why it exists. Given a raw schema, a competent model writes SQL that is *plausible*. It will use `AVG(net_revenue)` for average order value, which is wrong twice over — the tax and the average-of-averages. It will `INNER JOIN` the customer dimension and silently drop 5,664 orders. Neither error produces an exception; both produce a confident number that is wrong by an amount nobody notices. The semantic layer is what converts plausible SQL into correct SQL, and it is the single highest-leverage component in the system.

Entity resolution

`app/semantic/entities.py` resolves what a question *names* to what the database *stores*, deterministically and before the planner sees the question. It loads 211 distinct values across 16 dimensions at startup, then matches in decreasing order of certainty: exact, alias, identifier normalisation, word-boundary substring, and finally fuzzy matching at a 0.84 difflib threshold restricted to single tokens of five characters or more.

It handles `ST7` → `ST007`, `Bangalore` → `Bengaluru` (17 city aliases), `Bengalru` → `Bengaluru`, `veg burgers` → the three Veg Burger SKUs, and `delivery` → `Swiggy` and `Zomato` together. The canonical values and the columns they live in are injected into the planner prompt as facts.

Three properties matter more than recall, and each is tested:

- **It never invents.** `ST999` resolves to nothing, not to the nearest store.
- **It never widens.** `veg burger` is a substring of `Non-Veg Burger 2`; matching on a plain substring would answer a question about veg burgers with non-veg data, so matches must begin at a word boundary.
- **Ambiguity survives.** A phrase matching several *kinds* of entity — `pizzas` matches both the Pizza category and several product names — is reported as ambiguous rather than decided, which is what drives the clarification path.

6. Verification

Thirteen deterministic checks run on every request, in Python, before any model is consulted:

Check	Catches
<code>all_sub_queries_executed</code>	A planned sub-query that never ran
<code>results_non_empty</code>	A query that returned nothing
<code>no_all_null_columns</code>	A column that is null in every row
<code>no_negative_measures</code>	Negative revenue, order counts or quantities
<code>aov_reconciles</code>	AOV that is not $\text{revenue} \div \text{orders}$, within 0.5 INR
<code>parts_sum_to_total</code>	A breakdown that does not sum to its own stated total, within 1 INR
<code>shares_sum_to_100</code>	Percentage columns not summing to 100, within 1 point
<code>revenue_within_plausible_bound</code>	A figure above the configured plausibility ceiling
<code>row_count_matches_expectation</code>	"Top 5" returning four rows or six
<code>sql_references_planned_metrics</code>	Executed SQL that does not mention the metrics the plan named
<code>dates_within_data_range</code>	Any date outside 2025-08-01 to 2026-07-31
<code>results_cover_the_window</code>	A window silently truncated — see below
<code>dimension_coverage</code>	A breakdown missing members of its own dimension — see below

A model escalation (`llm_plausibility`) is consulted **only** when no error-severity check has failed and something remains genuinely ambiguous. Its verdict is capped at warning severity: it can raise a concern, never overturn arithmetic.

Why deterministic checks come first. Arithmetic is strong evidence and a model grading its own pipeline is weak evidence. Running the model first, or letting it adjudicate a failed sum, means it can sometimes talk the pipeline out of a real defect — the exact failure this ordering prevents. Verification that depends on a provider also disappears when the provider does; these checks cost nothing and work in a total outage.

Two checks written in response to specific defects

Both caught errors in which **every individual figure was correct**, which is precisely the class of error no amount of number-checking finds.

`results_cover_the_window`. A query filtered `month_key BETWEEN '2026-05-01' AND '2026-07-31'`. `month_key` is a 'YYYY-MM' string, and '2026-05' sorts before '2026-05-01', so every May row was silently dropped. The remaining June and July figures were individually correct, the totals were internally consistent, and the answer described a two-month window as three. The check now asserts that a result claiming to cover a window contains every month in it. A SQL rule was added alongside it stating that `month_key` is not a date.

`dimension_coverage`. A "by city" question returned one city. Nothing was arithmetically wrong; the answer simply described a fraction of the population as though it were the whole. The check compares distinct values returned against the dimension's known cardinality and flags a shortfall.

7. Reliability

The deployed URL may be opened months after the keys configured for it expired. Reliability here is therefore a correctness requirement, not an optimisation.

Multi-provider failover. Four providers — Anthropic, OpenAI, Gemini, xAI — behind one client, called in configured order. Each provider's HTTP API is called directly through `httpx` rather than through a vendor SDK; four SDKs would mean four dependency trees and four sets of abstractions to read past, while the raw request and response shapes are short enough to read in one sitting. Transient failures (429, 5xx, timeout) retry within the provider; permanent ones (401, 400) fail over immediately, because retrying a rejected key only wastes the seconds a user is waiting.

Circuit breaker. After 2 consecutive failures a provider is skipped for a 120-second cooldown, then tried again. Without this, every request pays the full timeout of a provider whose key expired months ago. Two rather than one, because a single failure is often the request's fault rather than the provider's. When the breaker has taken *everything* out of rotation it is bypassed and the providers are tried anyway: an open breaker is a prediction that a call will fail, and a prediction is not a good enough reason to return nothing.

Startup liveness probe. One minimal call per provider at startup, run as a background task so a slow provider cannot delay the port opening. Providers that answer are ordered ahead of those that did not, and a failed probe opens the breaker immediately — the probe *is* the evidence, and making the first user request rediscover it defeats the point. On the current deployment this correctly identified in about four seconds that three of four providers had no credit remaining.

A malformed JSON response fails over to a different provider. Malformed structure is a property of the model, not of the moment: at temperature zero the same model given the same prompt usually produces the same broken output, so retrying it burns a user's time to arrive back where it started. Each provider gets exactly one attempt per JSON call.

Answer caching. Answers are cached by normalised question text and the cache file is committed to the repository. The consequence is the property this system is built around:

*All eight evaluation questions answer completely — headline, narrative, executed SQL, verification report and four-agent trace — with **no LLM provider configured at all**.*

This is asserted by `scripts/check_offline.py`, which clears every provider key before importing the application, and which **runs in CI on every push**. It is a property enforced by a script rather than a claim made in a document.

Provider health is exposed. `GET /api/providers` reports which providers are configured, which are healthy, which are in cooldown and for how long, and each one's last probe result. `POST /api/providers/probe` re-probes on demand, so a provider whose key was fixed can be brought back without a redeploy. No credential material appears in either payload, not even a masked prefix.

8. Security

The system is publicly deployed and a language model writes its queries, so any text a user types reaches a SQL generator. Prompt injection is an expected input, not a hypothetical one.

Validation is parse-based, never pattern-based. Regex over SQL loses to comments, string literals, nested quoting and whitespace. Every check inspects the sqlglot AST. The difference is not theoretical:

- `SELECT store_name FROM dim_store WHERE store_name LIKE '%DROP TABLE%' LIMIT 5` is **accepted**. It contains the word `DROP` inside a string literal, which a keyword blocklist would reject and a parser correctly sees as a `SELECT` over an allowlisted table.
- `SELECT 1 -- ' \n; DROP TABLE dim_store` is **rejected**. It defeats a naive pattern by hiding the second statement behind a comment, but the parser sees two statements and the first layer refuses anything that is not exactly one.

Six layers, none load-bearing on its own.

- 1 The SQL must parse as exactly one SQLite statement.
- 2 That statement must be a query node (`SELECT` or a set operation), checked on the parsed node class, so `DROP` cannot be disguised.
- 3 Every table read must be on the semantic layer's ten-table allowlist, with CTEs defined in the same query resolved first so they are not mistaken for unknown tables.
- 4 No `sqlite_*` internal table may be touched.
- 5 A `LIMIT` is required, and injected when missing; results are capped at 1,000 rows and queries at 10 seconds.
- 6 Execution happens over a **read-only URI connection** (`file:...?mode=ro`).

Layer 6 is the one that matters most. Layers 1–5 are software and can have bugs; a connection opened `mode=ro` is enforced by SQLite itself. `tests/test_sql_guard.py` contains 75 tests, including injection attempts through the question text.

Rate limiting. A per-IP sliding window, 10 new analyses per minute and 200 per day. Cached answers are exempt: they cost nothing to serve, and rate-limiting them would punish exactly the interaction the system wants a reviewer to have.

Key handling. Keys are read from the environment, never logged, and never returned by any endpoint — health output reports which providers are *configured*, never their credentials. `tests/test_llm.py` and `tests/test_resilience.py` assert that no key-shaped string appears in a health payload, using deliberately key-shaped fixtures.

9. Testing

813 tests, all passing with no API keys configured.

Area	Tests	Area	Tests
Contracts	101	Analyst	37
Semantic layer	95	Insight	35
SQL guard	75	API	32
Golden answers	67	Verifier	29
Entity resolution	52	Resilience	29
ETL	50	Config	25
Golden end-to-end	50	Quality gate	24
LLM client	37	Planner	21
Frontend contract	18	Degraded paths	16
Orchestrator	16		

Ground truth is computed independently. The eight evaluation answers are computed from the Excel workbook with pandas and stored in `tests/golden_answers.json`. The application answers from the SQLite star schema. `tests/test_golden.py` runs the SQL equivalent of five of those questions and asserts it lands within 1 INR of the

pandas figure. Two independent implementations agreeing across twelve months of data is real evidence the ETL is faithful; either one alone proves nothing. The same file also checks internal consistency, so a mistake in any single answer surfaces as a contradiction rather than a plausible number.

The end-to-end suite (`tests/test_golden_e2e.py`, 39 assertions across 50 tests) drives the full pipeline and asserts, per question, the specific facts that matter — that Q5 states plainly that *no* city declined every month, that Q8 names exactly nine stores and separates the four above their own baseline from the five genuinely below it, that no store identifier is malformed. It also applies universal checks to every question: answered, verification not failed, all four agents present in the trace, and no figure in the narrative that does not appear in the query results.

Replay versus live. By default the suite replays the plan, SQL and narrative that the live system actually produced, then runs them through the real guard, the real SQLite database and the real verifier. So CI validates the data path — the part that must never break silently — without a network call or a cent of spend. Setting `QUICKBITE_E2E_LIVE=1` runs the same assertions against live providers.

The honest limitation of replay is that **it cannot catch a prompt regression**. If a prompt change causes the planner to produce a worse plan, replay will not notice, because it replays the old plan. Replay proves the data path is intact; only a live run proves the prompts still work. Both exist for that reason.

Environment pinning. The suite had no `conftest.py`, so it inherited whatever provider keys the developer happened to have in a gitignored `.env`. Twelve API tests silently depended on that: they POST to `/api/ask`, which short-circuits at the provider guard when nothing is configured, so the stubbed orchestrator was never reached. They passed locally and failed in CI — the same suite exercising different code paths in two environments. `tests/conftest.py` now pins the environment before `app.config` is imported: one obviously-fake provider key, the other three blank, and the startup probe disabled so no test can make a network call.

Frontend contract tests. The TypeScript types in `frontend/lib/types.ts` are hand-maintained mirrors of the Pydantic contracts, and nothing at runtime checks them — a field added on one side and forgotten on the other compiles cleanly and reads `undefined` in the client. `tests/test_frontend_contract.py` compares all fifteen wire contracts field by field, plus both string-union enums. It was written after exactly that drift occurred.

10. Engineering decisions

Why a custom orchestrator rather than a framework

LangChain or LlamaIndex would have supplied the agent loop. They would also have supplied their own retry semantics, their own error types, their own prompt templates and their own opinions about memory — and the whole argument of this system is that the *specific* degradation behaviour is the product. The self-healing retry that re-runs only the sub-queries a verification report blames, the distinction between a provider being unreachable and a model writing bad SQL, the trace that shows skipped steps rather than hiding them: each of these is a decision I wanted to make explicitly and be able to point at. The orchestrator is about 870 lines and every branch in it is deliberate. A framework would have made the first day faster and every subsequent day slower.

Why deterministic verification precedes model verification

Because arithmetic is strong evidence and a model reviewing its own pipeline is weak evidence. If the model runs first, or is allowed to adjudicate a failed check, it will sometimes talk the system out of a real defect. The escalation therefore runs only when nothing arithmetic has failed, and its verdict is hardcoded to warning severity. There is also a reliability argument: checks that depend on a provider vanish when the provider does, and these do not.

Why qualifying sets and comparisons are computed in SQL

This is the decision that most improved answer quality, and it was established by three separate defects.

Defect 1 — a misclassified store. Asked which stores were consistently declining and why, the system named a store as the top concern that was, in fact, still above its own prior-quarter baseline. Every figure quoted was correct; the store had fallen from an unusually strong quarter and was reverting, not deteriorating. The model had been given two result sets and asked to compare them by reading. The fix was structural: the baseline sub-query now returns one row per entity with `window_revenue`, `baseline_revenue`, `delta_abs`, `delta_pct` and `is_above_baseline` already computed as columns. The baseline is a **column, never a filter** — making it a filter causes the reverting entities to vanish from a question that asked about all of them.

Defect 2 — a headline contradicting its own findings. The answer to "are any cities declining" led with a city named as declining, while its own key findings listed the monthly figures showing that city rising in the middle month. The model had been handed 8 cities × 3 months and asked to spot a monotonic pattern by eye. The fix: one sub-query computes the qualifying set in SQL with window functions and emits a boolean column; the monthly series remains a separate sub-query so the answer can still show the shape of each decline.

Defect 3 — an evasive answer. After a related change, the same question produced "several of these locations", naming none. The correct answer was that *no* city qualified — and an empty set is a perfectly good answer, but only if the system computes the set rather than inferring it. A yes/no question needs the set just as much as a "which" question does.

The general rule this established: **never ask the model to derive a comparison it could read**. Fifty stores across three months is a hundred and fifty numbers, and a reader comparing them by eye will miss some. Compute the membership test in SQL, emit a boolean column, pre-group it in Python for the prompt.

Why instruction scoping lives in code, not in prompt conditions

The baseline-columns instruction above was first added as a conditional sentence in the shared analyst prompt: *"when this sub-query is the baseline comparison, return one row per entity..."*. It contaminated unrelated sub-queries. Two questions' monthly-trend queries became copies of the baseline query, collapsing a three-month series into two period totals — destroying exactly the thing a trend question asks about — and both answers then concluded nothing had declined.

Prose scoping does not reliably scope. The fix moved classification into Python: functions that inspect a sub-query's id and purpose and decide which instruction block it receives, checked in a defined order because the qualifying query for "declined every consecutive month" contains the words "consecutive" and "month" and would otherwise be classified as a trend query. Scoping an instruction is a control-flow decision, and control flow belongs in code.

Why schema examples are stripped from the injected response schema

The insight agent is given `Insight.model_json_schema()` so it knows what shape to return. That schema embeds the model's `json_schema_extra` example — and the example contained real findings from this dataset, written when the contract was authored. A live run reproduced "48,860 INR" and "ST039 fell 41.4%" verbatim in an answer to a question that had never queried city data. The model was not hallucinating; it was copying an example it had been handed and told was the correct shape.

`schema_without_examples()` now strips every `examples` and `json_schema_extra` key from the schema before injection. The general lesson: anything placed in a prompt is training data for that call, including the parts intended only as formatting guidance.

Why the accent colour avoids the hues that already carry meaning

The interface uses one accent colour. It was originally a deep pine green, which looked correct until the in-flight trace was examined: the running step's dot was nearly indistinguishable from the completed steps' dots, because "verified" and "succeeded" are also green. Worse, chart bars in accent green read as *good* rather than as *data* — a semantic claim the chart was not making.

Green, amber and red are spoken for by verification status. The accent moved to a deep aubergine, which collides with nothing, carries white text at 12:1 contrast, and lets a bar chart be neutral. Colour in this interface is a channel that already carries meaning, so the accent had to be chosen from what was left rather than from what looked nice.

Why answers are cached and committed

A reviewer may open the deployed URL long after the API keys have expired, and a system that cannot answer its own evaluation questions at that moment has failed regardless of how well it is built. The cache makes the eight questions a property of the repository rather than of a live credential. Serving is cache-first, so those answers are also instant, which is the right behaviour for the interaction a reviewer is most likely to have. The cost is that a cached answer is a frozen snapshot: the committed reports were computed when each answer was warmed, so six of the eight show twelve verification checks rather than the thirteen the current code runs.

11. Known limitations, and what I would do next

Stated plainly, because a reviewer should finish this section believing the author understands the system's weaknesses.

Narrative referential slips pass verification. Asked to compare two stores, one live answer said "ST015 outperformed ST015" where the second should have been ST007. Every *number* was correct, so all thirteen checks passed. The verifier checks arithmetic and traceability of figures; it does not check that pronouns and entity references in prose are internally consistent. A check comparing entity mentions in the narrative against the entities present in the result rows would catch this, and is the next check I would write.

Insight degrades under provider rate limits. In live testing, two of nineteen questions lost their narrative to a free-tier rate limit and fell back to deterministic prose. The data and verification were intact and the trace said what happened, but the answer was noticeably poorer. This is a funding constraint rather than a design flaw — only one provider has credit — but the effect on a reviewer is the same.

A question naming a non-existent entity returns empty results rather than saying so. The resolver correctly declines to invent a match for ST999, but nothing downstream converts "the question named an identifier and it resolved to nothing" into a clear message. The user gets an honest but unhelpful "no data could be retrieved". The fix is small and specific: when a question contains an identifier-shaped token that resolves to nothing, say that the entity does not exist before planning.

The deterministic SQL fallback covers only single-table queries. When no provider is reachable, SQL can be assembled from the plan — but only over `fact_orders`, and only for metrics and dimensions available there. Anything needing a join is refused rather than guessed, which is the correct trade (a wrong join changes the grain silently) but means margin, product and city questions have no degraded path.

Clarification is conservative, and its false-negative rate was not measured. The planner asks for clarification only where guessing would mislead, and defaults everywhere else. I verified it triggers on two genuinely ambiguous questions and that the round trip works. I did **not** sweep for questions that *should* have been clarified and were silently answered under an assumption instead. That rate is unknown.

Coverage is estimated, not measured. Based on nineteen live questions across the question types the system claims to support, I would estimate 80–85% of reasonable questions about this dataset now answer correctly. That

is an estimate from a small sample, not a benchmark. A proper evaluation set of a hundred questions with human-labelled expected answers is what this needs, and is the single thing I would build next.

Other known gaps. Year-on-year comparison is impossible and correctly refused, but the dataset covering only twelve months limits every seasonality question to a single observation of each season. The chart type is chosen by the insight agent from a closed set and is occasionally a poor fit. There is no authentication; the deployment is public by design and rate-limited rather than gated.